



Almost Continuous Integration for the Co-Development of Highly Integrated Applications and Third Party Libraries

Roscoe A. Bartlett

<http://www.cs.sandia.gov/~rabartl/>

Department of Optimization & Uncertainty Estimation
Trilinos Software Engineering Technologies and Integration Lead

Sandia National Laboratories

Sandia Software Engineering Seminar Series
January 14, 2009



Applications (APPs) and Third-Party Libraries (TPL) at SNL

Applications (APPs)	Third Partly Libraries (TPL)
Aleph	Trilinos
Xyce	Trilinos
Titan (VTK)	Trilinos
Charon*	Trilinos*, Xyce, Nevada
Alegra*	Trilinos*, Xyce, Nevada
SIERRA*	Trilinos*

* Some experimentation with more frequent APP + TPL Integration

- Tighter level of APP + TPL integration is needed in many cases
- Co-development of APP + TPL(s) is often needed to drive new efforts
- Current software engineering infrastructure and practices are insufficient to support desired goals
- We need new software engineering infrastructure to support these integration efforts




Lean/Agile Software Engineering Principles

- High quality software is developed in small increments and with sufficient testing in between sets of changes.
- High quality defect-free software is most effectively developed by not putting defects into the software in the first place (i.e. TDD, code reviews, pair programming, etc.).
- Software should be delivered to real (or as real as we can make them) customers in short intervals.
- Ruthlessly remove duplication in all areas.
- Avoid points of synchronization. Allow people to work as independently as possible and have the system set up to automatically support this.
- Most mistakes that people make are due to a faulty process/system (W. Edwards Deming).
- Automation is needed to avoid mistakes and improve software quality.

References: <http://www.cs.sandia.gov/~rabarl/readingList.html>



Lean/Agile Methods: Development Stability




Problems

- Cost of fixing defects increases the longer they exist in the code
- Difficult to sustain development productivity
- Broken code begets broken code (i.e. broken window phenomenon)
- Long time between branch and release
 - Difficult to merge changes back into main development branch
 - Temptation to add “features” to the release branch before a release
- High risk of creating a regression



Lean/Agile Methods: Development Stability




Advantages

- Defects are kept out of the code in the first place
- Code is kept in a near releasable state at all times
- Shorten time needed to put out a release
- Allow for more frequent releases
- Reduce risk of creating regressions
- Decrease overall development cost




APP Only Upgrades After Each Major Release of TPL



- Transition from TPL X to TPL X+1 can be difficult and open ended
- Large batches of changes between integrations
- Greater risk of experiencing real regressions
- Upgrades may need to be completely abandoned in extreme cases
- However, this is satisfactory for many APP+TPL efforts!




Build and Test APP Against both TPL Release and TPL Dev



- APP (SIERRA) Dev Developers only build/test against TPL Release
- TPL (Trilinos) Dev Developers work independent from APP
- Changes between TPL Release and TPL Dev handed through a) Refactoring, b) minimal ifdefs (NO BRANCHES)! => **Backward Compatibility!**
- Use of staggered releases of TPL and APP
- APP + TPL Dev Developers drive new capabilities
- Difficult for APP to depend too much on TPL
- Does not support tighter levels of integration
- However, this is satisfactory for many APP+TPL efforts!



APP Dev Builds Against Both TPL Release and TPL Dev



- All changes are tested in small batches
- Low probability of experiencing a regression
- Extra computing resources to test against 2 (3) versions of TPL
- Some difficulty flagging regressions of APP + TPL Dev
- APP developers often break APP + TPL Dev
- Difficult for APP to rely on TPL too much
- Hard to verify TPL for APP before APP release
- However, this is satisfactory for many APP+TPL efforts!



APP + TPL Integration: Different Collaboration Models

- APP Dev only upgraded after each major release of TPL
 - Little to no testing of APP + TPL Dev in between TPL releases
- APP Dev builds against both TPL Release and TPL Dev
 - APP developers work against TPL Release
 - APP + TPL team(s) build against TPL Dev
 - Daily integration testing done for both APP + TPL Release and Dev
 - Staggered releases of TPL and APP
- APP Dev developed only against TPL Dev (with “Almost” Continuous Integration)
 - Regular APP developers work independently using very recent APP-owned VC copy of TPL Dev-
 - Regular TPL developers work independently
 - APP Dev + TPL Dev developers
 - Check-out and modify APP Dev
 - Check-out and modify TPL Dev
 - Run both APP and TPL pre-checkin test suites
 - Check into both APP-owned and main TPL VC repositories
 - Nightly testing of APP Dev + TPL Dev automatically updates APP-owned TPL Dev- VC Repository
 - Releases best handled as combined releases of APP and TPL




General Development & Testing Principles

- Regular TPL developers only build and run TPL pre-checkin test suite.
- Regular APP developers should only check out code that has already built and passed the pre-checkin APP test suite.
- Nightly APP regression (and other) tests should only be run on code that has already been shown to build and pass the pre-checkin APP test suite.
- Code that builds and passes the pre-checkin test suite is safe to check in.




Basic Setup for APP + TPL Almost Continuous Integration






Standard APP Development Process



- TPL (Dev-) code is typically not modified by average APP developers!
- However, small changes can be made and can be good!




APP Dev + TPL Dev Development Process



- Pre-checkin test suites for APP and TPL are both run before checkin
- Simultaneous checks into APP-owned TPL Dev- and Main TPL Dev VC Repositories!
 - Changes in APP-owned TPL VC Dev- Repos get back into Main TPL VC Dev Repos!



Nightly APP + TPL Dev Testing and Checkins of TPL Dev-



- Only runs pre-checkin test suite and then only on the primary development platform! (just like a regular APP developer)
- TPL Dev- VC Repository is automatically updated by nightly testing process if a) merge, b) build, and c) pre-checkin test suite all pass!
 - This is the same criteria we have for any regular APP developer checkin!
- Integration build is checked throughout the day with continuous integration (but without the auto-updates of TPL Dev- VC repository to avoid conflicts)




APP + TPL Development and Testing Details and Policies


- Nightly Testing:
 - Nightly APP Dev + TPL Dev testing and checking in only run on primary development platform and only runs pre-checkin test suite
=> Minimizes extra testing computer resources!
 - Nightly APP regression (and other stronger) tests are only run on APP Dev + TPL Dev- and *not* with TPL Dev (but on the same day after upgrade of APP Dev-)
 - Only one version of Dev code goes through extended testing (e.g. porting, regression, performance, scalability).
 - If APP Dev + TPL Dev testing and updating of TPL Dev- succeeds, then extended testing will involve all changes to APP Dev and TPL Dev in the last 24 hours.
- Continuous Integration Testing:
 - Build and test APP Dev + TPL Dev throughout the day to flag problems and to help support co-development of APP Dev + TPL Dev
- Open Questions:
 - How are multiple TPL handled in nightly testing ?
 - Are all TPL updated at the same time in nightly testing process?
 - Are TPL updated and testing separately in a chain (TPL 1 followed by TPL 2, etc.)?
 - What about intra-TPL dependencies (i.e. Nevada and Xyce => Trilinos)?
 - Do all TPL need to follow this process as well?



APP + TPL Almost Continuous Integration and Releases



The Future of APP + TPL Integration?



Nightly Testing: APP Dev + TPL Dev (pre-checkin tests only, TPL Dev- checkin)


Nightly Testing: APP Dev + TPL Dev- (complete test suites)

Supported with asynchronous continuous integration testing of APP Dev + TPL Dev

- All changes are tested in small batches
- Low probability of experiencing a regression between major releases
- Less computing resources for detailed nightly testing (only one TPL version)
- All tested regressions are flagged in less than 24 hours
- Allows code to flow freely between the APP and TPL
- Supports rapid development of new capabilities from top to bottom
- All code checked out by APP Dev developers has passed pre-checkin build/test
- More complex processes (i.e. requires some tools?)
- APP Dev developers spend more time spent recompiling TPL code
- Recommended for projects requiring high levels of integration & collaboration!



Challenges with APP-Specific TPL Releases



Multiple releases of TPL (Trilinos) presents a possible problem with complex APPs

Solution:

=> Provide perfect backward compatibility of Trilinos (TPL) X through Trilinos SIERRA Y+1




Assorted Ideas for APP Dev + TPL Dev Nightly Testing

- Nightly and continuous updating, testing, and checkin algorithm
 - Check out APP Dev and + TPL Dev- from APP-owned TPL Dev- VC Repository(s)
 - Build and run pre-checkin APP test suite (for APP Dev + TPL Dev-)
 - For each TPL ($i = 0 \dots N-1$) [In order of increasing dependencies]
 - Perform update of TPL i Dev from main TPL i VC Dev repository
 - Build and run pre-checkin APP test suite
 - If all passed, check into APP-owned TPL i Dev- VC repository [Nightly only]
 - Otherwise, skip checkin into APP-owned TPL i Dev- VC repository
- Advantages
 - Failures with one TPL do not automatically bring down integration with all TPL
 - Example: If Trilinos Dev works with Charon but Xyce Dev does not, at least Trilinos Dev would get updated and used by Charon Dev.
 - Provides additional information on where regressions are coming from
 - Example: A test passes with APP Dev + TPL Dev- but fails with APP Dev + TPL Dev



Maintenance of APP + TPL Integration



- **APP + TPL Monitor:**

- Member of the APP development team
- Has good familiarity with the TPLs
- Performs first-round triage (APP or TPL?)
- Forwards issues to APP or TPL Reps
- Ultimate responsibility to make sure issues are resolved

- **APP Representative:**

- Member of the APP development team
- Second-round triage of APP issues
- Forwards hard APP issues to APP developers

- **TPL Representative:**

- Member of the TPL development team
- Has some familiarity with the APPs
- Second-round triage for TPL issues
- Forwards hard TPL issues to TPL developers

- **General principles:**

- Roles of authority and accountability (Ordained by management)
- At least two people serve in each role
- Rotate people in roles



Summary #1

- Nightly building and testing of the development versions of the application and TPLs:
 - results in better production capabilities and better research,
 - brings TPL developers and APP developers closer together allowing for a better exchange of ideas and concerns,
 - refocuses TPL developers on customer efforts,
 - helps drive continued research-quality TPL development, and
 - reduces barriers for new TPL algorithms to have impact on production applications.
- APP Dev developed only against TPL Dev (with “Almost” Continuous Integration)
 - Regular APP developers work independently using very recent APP-owned VC copy of TPL Dev-
 - Regular TPL developers work independently
 - APP Dev + TPL Dev developers
 - Check-out and modify APP Dev
 - Check-out and modify TPL Dev
 - Run both APP and TPL pre-checkin test suites
 - Check into both APP-owned and main TPL VC repositories
 - Nightly testing of APP Dev + TPL Dev automatically updates APP-owned TPL Dev- VC Repository
 - Releases best handled as combined releases of APP and TPL



Summary #2

- Integration Models:

- APP Dev only upgraded after each major release of TPL
 - Little to no testing of APP + TPL Dev in between TPL releases
- APP Dev builds against both TPL Release and TPL Dev
 - Daily Integration testing done for both APP + TPL Release and Dev
 - Staggered releases of TPL and APP
- APP Dev developed only against TPL Dev (with “Almost” Continuous Integration)
 - APP Dev + TPL Dev developers update both APP-owned and main TPL repositories
 - Nightly testing of APP Dev + TPL Dev automatically updates APP-owned TPL Dev- VC Repository
 - Releases best handled as combined releases of APP and TPL
 - TPL Dev- checkins can be dialed back approaching TPL Release and Dev Integration!

- Final thoughts

- Each of these different integration models will be appropriate for a particular APP+TPL situation.
- The particular integration model can be switched during the life-cycles of the APP and TPL depending on several factors:
 - How critical is the TPL functionality currently to the APP?
 - Are there alternatives to a particular TPL that can duplicate functionality?
 - How actively is the TPL being developed?
 - Is it critical for the APP to continue to accept new releases of the TPL?
 - How active is the collaboration between APP and TPL developers?
 - Is the TPL a fundamental part of the infrastructure of the APP?
 - ...